



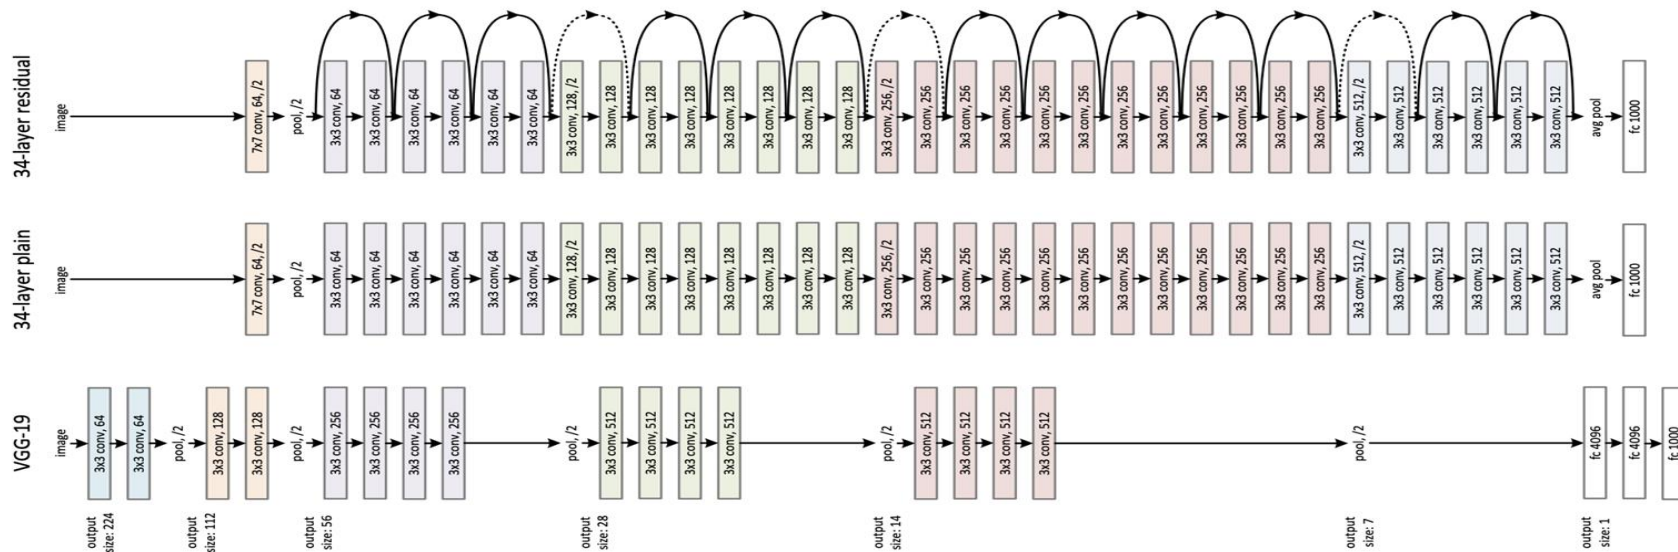
Residual Network Model (RNN) Exploration and Implementation

Presented by Yiru Xiong



ResNet Recap

- A Convolutional Neural Network (CNN) based deep learning model
- Overcome Vanishing Gradient problem when increase the number of layers
- Skip connections connect activation of a layer to further layer while skipping some in between
- Stack residual blocks together
- Generally used for computer vision applications - image classifications, etc.



Data

- The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class.
- There are 50000 training images and 10000 test images.
- Classes are completely mutually exclusive.

```
# CIFAR-10 dataset
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

batch_size = 5

trainset = torchvision.datasets.CIFAR10(root='/content/drive/MyDrive/data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='/content/drive/MyDrive/data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

airplane



automobile



bird



cat



deer



dog



frog



horse




ship



truck



Baseline Model



- A simple CNN model

```
=====
Layer (type:depth-idx)                Param #
=====
CNN                                     --
|---Conv2d: 1-1                         456
|---MaxPool2d: 1-2                      --
|---Conv2d: 1-3                       2,416
|---Linear: 1-4                       48,120
|---Linear: 1-5                       10,164
|---Linear: 1-6                        850
=====
Total params: 62,006
Trainable params: 62,006
Non-trainable params: 0
=====
```

- Optimizer: SGD
- Learning rate: 0.001
- Momentum: 0.9
- Epoch_num: 10

**Accuracy on
test images:
62%**

ResNet

- Optimizer: SGD
- Learning rate: 0.001
- Epoch_num: 10
- With learning rate decay

Accuracy on test images: 80.95%

Layer (type:depth-idx)	Param #
ResNet	--
└─Conv2d: 1-1	432
└─BatchNorm2d: 1-2	32
└─ReLU: 1-3	--
└─Sequential: 1-4	--
└─ResidualBlock: 2-1	--
└─Conv2d: 3-1	2,304
└─BatchNorm2d: 3-2	32
└─ReLU: 3-3	--
└─Conv2d: 3-4	2,304
└─BatchNorm2d: 3-5	32
└─ResidualBlock: 2-2	--
└─Conv2d: 3-6	2,304
└─BatchNorm2d: 3-7	32
└─ReLU: 3-8	--
└─Conv2d: 3-9	2,304
└─BatchNorm2d: 3-10	32
└─Sequential: 1-5	--
└─ResidualBlock: 2-3	--
└─Conv2d: 3-11	4,608
└─BatchNorm2d: 3-12	64
└─ReLU: 3-13	--
└─Conv2d: 3-14	9,216
└─BatchNorm2d: 3-15	64
└─Sequential: 3-16	4,672
└─ResidualBlock: 2-4	--
└─Conv2d: 3-17	9,216
└─BatchNorm2d: 3-18	64
└─ReLU: 3-19	--
└─Conv2d: 3-20	9,216
└─BatchNorm2d: 3-21	64
└─Sequential: 1-6	--
└─ResidualBlock: 2-5	--
└─Conv2d: 3-22	18,432
└─BatchNorm2d: 3-23	128
└─ReLU: 3-24	--
└─Conv2d: 3-25	36,864
└─BatchNorm2d: 3-26	128
└─Sequential: 3-27	18,560
└─ResidualBlock: 2-6	--
└─Conv2d: 3-28	36,864
└─BatchNorm2d: 3-29	128
└─ReLU: 3-30	--
└─Conv2d: 3-31	36,864
└─BatchNorm2d: 3-32	128
└─AvgPool2d: 1-7	--
└─Linear: 1-8	650
Total params: 195,738	
Trainable params: 195,738	
Non-trainable params: 0	

ResNet - Hyperparameter Tuning

Grid Search -

- Batch Size - 128
- Number of epochs - 20
- Learning rate - 0.001
- Optimization algorithm - Adam
- With/ without learning rate decav -with

**Accuracy on test images:
86.18%**

```
# optimization algorithm
optim_grid = [optim.SGD, optim.RMSprop, optim.Adagrad, optim.Adadelta, optim.Adam, optim.Adamax, optim.NAdam]
lr_grid = [0.001, 0.005, 0.007, 0.01, 0.05]
epochs_grid = [10, 30, 50, 80]
batch_grid = [100, 150, 200, 500]
```

ResNet - Continued

Adding dropout layer

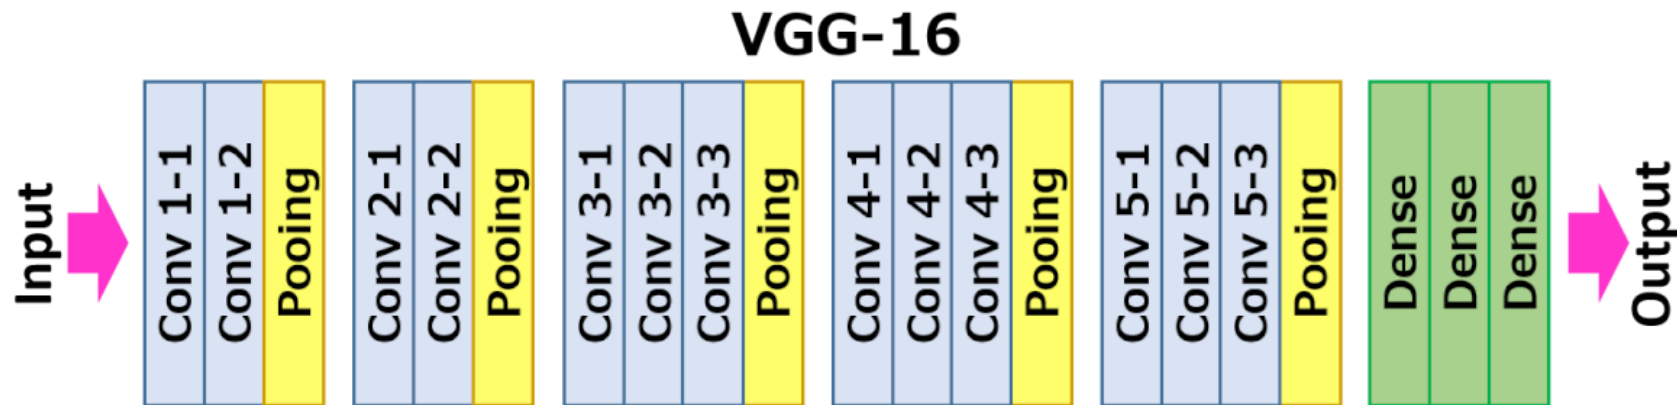
- Dropout layer can be added to visible layer
- Dropout layer can be added to hidden layers
- Dropout rate
0.2 - meaning one in five inputs will be
randomly excluded from each update cycle
- Applied after the activation layer

Accuracy on test images: avg 86%

Layer (type:depth-idx)	Param #
ResNet	--
└Conv2d: 1-1	432
└BatchNorm2d: 1-2	32
└ReLU: 1-3	--
└Sequential: 1-4	--
└ResidualBlock: 2-1	--
└Conv2d: 3-1	2,304
└BatchNorm2d: 3-2	32
└ReLU: 3-3	--
└Conv2d: 3-4	2,304
└BatchNorm2d: 3-5	32
└ResidualBlock: 2-2	--
└Conv2d: 3-6	2,304
└BatchNorm2d: 3-7	32
└ReLU: 3-8	--
└Conv2d: 3-9	2,304
└BatchNorm2d: 3-10	32
└Sequential: 1-5	--
└ResidualBlock: 2-3	--
└Conv2d: 3-11	4,608
└BatchNorm2d: 3-12	64
└ReLU: 3-13	--
└Conv2d: 3-14	9,216
└BatchNorm2d: 3-15	64
└Sequential: 3-16	4,672
└ResidualBlock: 2-4	--
└Conv2d: 3-17	9,216
└BatchNorm2d: 3-18	64
└ReLU: 3-19	--
└Conv2d: 3-20	9,216
└BatchNorm2d: 3-21	64
└Sequential: 1-6	--
└ResidualBlock: 2-5	--
└Conv2d: 3-22	18,432
└BatchNorm2d: 3-23	128
└ReLU: 3-24	--
└Conv2d: 3-25	36,864
└BatchNorm2d: 3-26	128
└Sequential: 3-27	18,560
└ResidualBlock: 2-6	--
└Conv2d: 3-28	36,864
└BatchNorm2d: 3-29	128
└ReLU: 3-30	--
└Conv2d: 3-31	36,864
└BatchNorm2d: 3-32	128
└AvgPool2d: 1-7	--
└Dropout: 1-8	--
└Linear: 1-9	650
Total params: 195,738	
Trainable params: 195,738	
Non-trainable params: 0	

Transfer Learning

- Leveraging pre-trained high-performance model
- VGG-16: Very Deep Convolutional Network for Large-scale Image Classification trained on ImageNet
- Convolutional Layers = 13, Pooling Layers = 5, Dense Layers = 3



Transfer Learning - continued


- VGG-16
- VGG-16 with batch normalization
- Experiment with different learning rate 0.001, 0.003

Num_epochs: 10, 20

Optimization Algorithm: Adam, RMSprop


- With the same parameter settings, adding batch normalization boosts the overall accuracy
- **Best Results:**
VGG-16: Accuracy on Testing Images: 89.93
VGG-16 with batch normalization: Accuracy on testing images; 91.03%

Experiment Results



Future Work

1. Vit - Vision Transformer. Attention Neural Network Architecture



2. Automatic Hyperparameter Tuning - Optuna



References

1. Tam, A. (2023, April 7). How to grid search hyperparameters for PyTorch Models. MachineLearningMastery.com.<https://machinelearningmastery.com/how-to-grid-search-hyperparameters-for-pytorch-models/>
 1. GitHub,github.com/rasbt/deeplearning-models/blob/master/pytorch_ipynb/transfer/transferlearning-vgg16-cifar10-1.ipynb. Accessed 21 Nov. 2023.
 1. “Accelerate Your Hyperparameter Optimization with PyTorch’s Ecosystem Tools.” Medium, 14 Sept. 2020, medium.com/pytorch/accelerate-your-hyperparameter-optimization-with-pytorchs-ecosystem-tools-bc17001b9a49.
- 